khaykingleb@gmail.com
/khaykingleb
/khaykingleb

GLEB KHAYKIN
Tbilisi, Georgia

+995 585-889-511
khaykingleb.com
@khaykingleb_blog

SUMMARY

MLOps/DevOps Engineer with 3.5+ years of software engineering experience. Specialized in boosting model performance, reducing costs by >$10k monthly, deploying scalable solutions for apps with >30k DAU, and optimizing data and workflow pipelines. Integrated state-of-the-art models and deployed infrastructure **across cloud** and **on-premise** environments. Background in **infrastructure**, **speech**, **natural language processing**, and **finance**.

EXPERIENCE

MLOps/DevOps Engineer at iClerk — San Francisco, USA (Remote)
Silicon Valley startup focusing on automating tedious tasks for B2B, $3.4M pre-seed investment — *July 2023 — Present*

○ Refined an ensemble model for **speech-to-text** and **speaker diarization**, achieving an **11x speedup** and improved accuracy
○ Developed a scalable inference pipeline with NVIDIA Triton, **reducing costs by 5x** for processing one-hour meetings
○ **Secured one of the 1st clients** by implementing **advanced RAG** for efficient parsing of financial statements
○ Architected a **K3s GPU cluster** with NVIDIA GPU Operator for monitoring and validation, **cutting GPU costs by >50%**
○ Established **IaC monitoring** & **alerting** for **4 K8s clusters** with Grafana Alloy OTel Collector, reducing downtime and errors

ML/MLOps Engineer at MynaLabs — Tbilisi, Georgia
Generative AI startup with $150M and $200M B2C app exits, both sold to Snapchat — *November 2022 — March 2024*

○ Built a REST/gRPC API for speech tasks with **priority queuing**, handling **60k+ inference requests** per hour
○ Researched and deployed state-of-the-art multilingual, emotional **text-to-speech** and **voice conversion** models **in 44 kHz**
○ Created a service for fast **speech model fine-tuning** to any voice in **under 1 minute**
○ **Terraformized** and migrated speech infrastructure to **Kubernetes** with **>4000 lines of IaC**
○ Introduced MLOps/DevOps **best practices to** a **20+** developer engineering **division**

Data Scientist at DiviAI — San Francisco, USA (Remote)
Silicon Valley startup building a financial data aggregator for the US market — *September 2021 — July 2022*

○ Engineered a PySpark pipeline for scraping S&P 500 news to estimate investor **sentiment with BERT**-like models
○ Developed **ETL pipelines** that aggregate data **from 150+** financial API **endpoints**
○ Designed a **listwise** learning-to-rank system for stocks **recommendation**, enhancing long-short strategy performance
○ Adopted a **Glicko rating** system to rank Russell 3000 CEOs

EDUCATION

NRU "Higher School of Economics" — Moscow, Russia
B.S. with summa cum laude in Economics and Data Science — *September 2018 — July 2022*

○ Relevant coursework: C++, Python, R, Machine Learning 1, Machine Learning 2, Large Scale Machine Learning, Deep Learning, Deep Learning in Audio Processing, Reinforcement Learning, Calculus, Linear Algebra, Probability Theory, Mathematical Statistics, Stochastic Processes, Econometrics, Microeconometrics, Differential & Difference Equations

PROJECTS

○ **DL for Audio**: Implemented and optimized models from scratch (e.g., **QuartzNet**, **FastSpeech**, **HiFi-GAN**)
○ **Huffman Archiver**: Built an ASCII archiver in **C++** with a **Conan**-based auto-build system
○ **Stochastic Optimization**: Constructed and tested **Particle Swarm Optimization** and **Genetic Algorithm**

MISCELLANEOUS

○ **CFA Level 1**: Passed in February 2021, scoring in the **top 10%** of candidates **worldwide**
○ **Teaching Assistant**: Facilitated course coordination in Accounting, **Probability**, **Statistics**, **Machine Learning** (2020–2022)

SKILLS

○ **Languages**: Python, Rust, Go, C++, SQL, TypeScript, JavaScript
○ **Frontend**: Remix, React, Tailwind, DaisyUI, Vite, Vercel
○ **Backend**: gRPC, REST, Websocket, Webhook, AsyncIO, FastAPI, Pydantic, JWT, Node.js
○ **DevOps**: Terraform, Ansible, Docker, Kubernetes, K3s, AWS, GCP, Grafana, Prometheus, Loki, Nginx, Acme, ArgoCD, CI/CD
○ **MLOps**: DVC, NVIDIA Triton, Ollama, LangChain, TorchScript, ONNX, TensorRT, CoreML, Gradio
○ **DE**: Hadoop, Spark, Databricks, Polars, Kafka, Celery, RabbitMQ
○ **Storage**: Redis, PostgreSQL, PGVector, Sqitch, S3, R2, Supabase, Airtable
○ **R&D**: PyTorch, Lightning, Hydra, W&B, Scikit-learn, Numpy, Pandas, SciPy